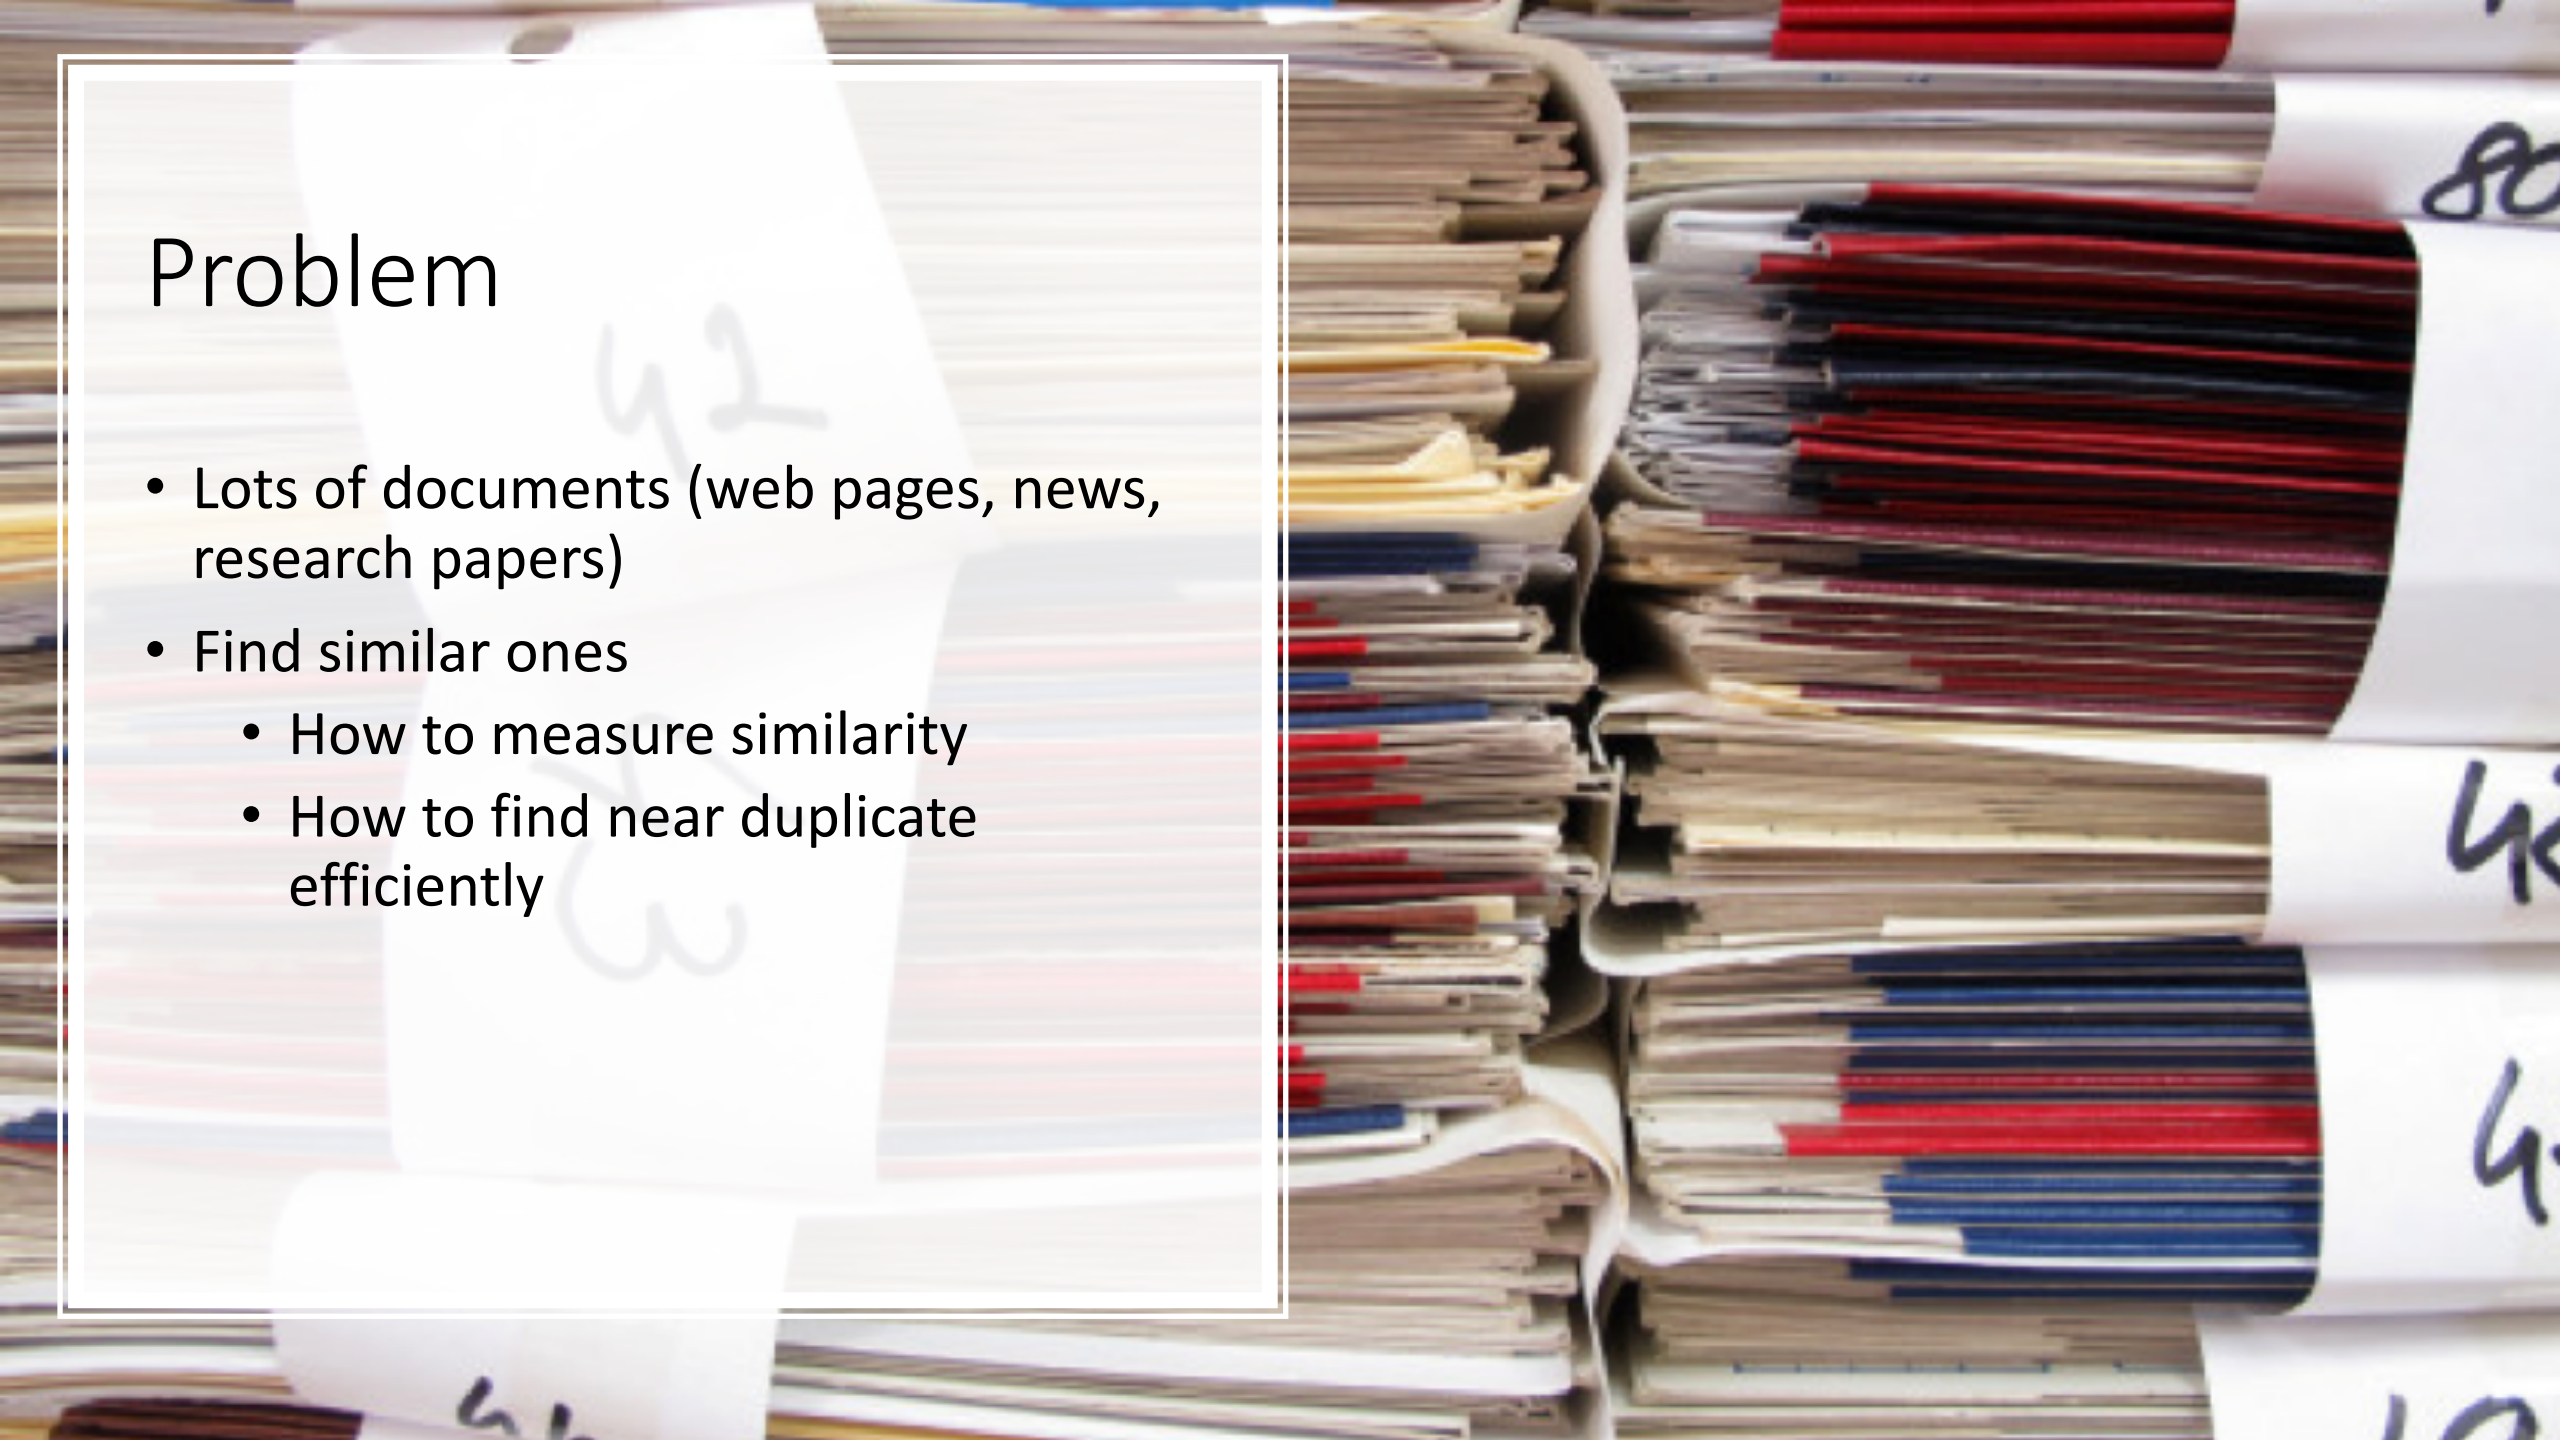


# Near Duplicate Detection using Simhash

Presented by  
Sumon Biswas

# Problem

- Lots of documents (web pages, news, research papers)
- Find similar ones
  - How to measure similarity
  - How to find near duplicate efficiently



# Applications

- Documents clustering
- Web mining
- Focused crawling
- Recommendation system
- Plagiarism checking
- Spam detection
- Computer vision

# Outline

- Similarity measure
- Locality sensitive hashing
- Simhash algorithm
- Hamming distance problem

# Similarity Measure

Two popular similarity measure between datasets:

Jaccard similarity:

$$\text{sim}(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

Cosine similarity:

$$\text{sim}(A, B) = \frac{|A \cap B|}{\sqrt{|A| |B|}}$$

# Near Duplicate Detection

- Exact duplicate detection is relatively easy
  - Using checksum technique
- Near duplicate is more challenging
  - find all web pages that are at least 90% similar to web page  $u$  (e.g., more than 90% words are same)
- To search  $O(n)$  comparisons are required
- To find all similar pairs in a collection  $O(n^2)$  comparisons are required

# Locality Sensitive Hashing

DEFINITION 1. *A locality sensitive hashing scheme is a distribution on a family  $\mathcal{F}$  of hash functions operating on a collection of objects, such that for two objects  $x, y$ ,*

$$\Pr_{h \in \mathcal{F}}[h(x) = h(y)] = \text{sim}(x, y) \quad (1)$$

# Simhash

- Documents are converted to set of features associated with weight
  - Features: word, shingle
  - Weight: frequency, TF-IDF
- Create  $b$  bit fingerprint of each document (e.g.,  $b = 64$ )
- To compute similarity between two documents compute hamming distance of two documents



---

**Algorithm 1** Simhash ( $u$ )

```

1:  $W \leftarrow$  array of  $b$  zeros
2: for  $i \in \mathcal{F}(u)$  do ▷ Examine each feature
3:    $\phi_i \leftarrow \text{UniformHash}(i)$  ▷ Compute  $b$ -bit hash
4:   for  $j = 1$  to  $b$  do ▷ Iterate through each bit
5:     if  $\phi_{ij} = 1$  then ▷  $j$ -th bit of  $\phi_i$ 
6:        $W[j] \leftarrow W[j] + w_i$  ▷ Add feature weight
7:     else
8:        $W[j] \leftarrow W[j] - w_i$  ▷ Subtract feature weight
9:     end if
10:  end for
11: end for
12: for  $j = 1$  to  $b$  do ▷ Revisit all bits
13:   if  $W[j] \geq 0$  then
14:      $B[j] \leftarrow 1$  ▷ Positive weight, set bit to 1
15:   else
16:      $B[j] \leftarrow 0$  ▷ Negative weight, set bit to 0
17:   end if
18: end for
19: return array  $B[1 \dots b]$  ▷ simhash

```

Tropical fish include fish found in tropical environments around the world, including both freshwater and salt water species.

(a) Original text

tropical 2 fish 2 include 1 found 1 environments 1 around 1 world 1  
including 1 both 1 freshwater 1 salt 1 water 1 species 1

(b) Words with weights

tropical	01100001	fish	10101011	include	11100110
found	00011110	environments	00101101	around	10001011
world	00101010	including	11000000	both	10101110
freshwater	00111111	salt	10110101	water	00100101
species	11101110				

(c) ~~8 bit hash values~~

1 -5 9 -9 3 1 3 3

(d) Vector  $V$  formed by summing weights

1 0 1 0 1 1 1 1

(e) 8-bit fingerprint formed from  $V$

# How it works

Given a collection of vectors in  $R^d$ , we define a LSH: We choose a random vector  $\vec{r}$ . Corresponding to this vector  $\vec{r}$ , we define a hash function  $h_{\vec{r}}$ :


$$h_{\vec{r}} = \begin{cases} 1 & \text{if } \vec{r} \cdot \vec{u} \geq 0 \\ 0 & \text{if } \vec{r} \cdot \vec{u} < 0 \end{cases}$$

Then for vectors  $\vec{u}$  and  $\vec{v}$ ,

$$\mathbf{Pr}[h_{\vec{r}}(\vec{u}) = h_{\vec{r}}(\vec{v})] = 1 - \frac{\theta(\vec{u}, \vec{v})}{\pi}$$

$\theta(\vec{u}, \vec{v})$  refers to the angle between vectors  $\vec{u}$  and  $\vec{v}$ . The function  $1 - \frac{\theta(\vec{u}, \vec{v})}{\pi}$  is closely related to the function  $1 - \cos(\theta)$ . In fact it is always within a factor 0.878 from it.

# Correlation



# Hamming Distance Problem

- If  $x$  and  $y$  are two binary numbers of length  $b$ 
  - Hamming distance  $H(x, y) = |\{i : x_i \neq y_i\}|$   
 $= |\{i : x_i \oplus y_i = 1\}|$
  - Similarity  $S(x, y) = 1 - H(x, y)/b$
- If we have 10,000 documents. Then we need about 50 million comparisons.

# Hamming Distance Problem

Index	Decimal	Hash(binary)	
1	37586	1001001011010010	
2	50086	1100001110100110	7
3	2648	0000101001011000	11
4	934	0000001110100110	9
5	40957	1001111111111101	9
6	2650	0000101001011010	9
7	64475	1111101111011011	7
8	40955	1001111111111011	4

Sort

Index	Decimal	Hash(binary)	
4	934	0000001110100110	
3	2648	0000101001011000	9
6	2650	0000101001011010	1
1	37586	1001001011010010	5
8	40955	1001111111111011	6
5	40957	1001111111111101	2
2	50086	1100001110100110	9
7	64475	1111101111011011	9

$H(3, 6) = 1$ ,  $H(8, 5) = 2$  are adjacent  
However,  $H(4, 2) = 2$  which has fall apart.

# Hamming Distance Problem

Rotate bits  
left twice



Index	Decimal	Hash(binary)	
4	3736	0000111010011000	
3	10592	0010100101100000	9
6	10600	0010100101101000	1
1	19274	0100101101001010	5
8	32750	0111111111101110	6
5	32758	0111111111101110	2
2	3739	0000111010011011	9
7	61295	1110111101101111	9

Sort




Index	Decimal	Hash(binary)	
4	934	0000001110100110	
2	2648	0000111010011011	2
3	2650	0010100101100000	11
6	37586	0010100101101000	1
1	40955	0100101101001010	5
8	40957	0111111111101110	6
5	50086	0111111111101110	2
7	64475	1110111101101111	6

Now,  $H(4, 2) = 2$  are adjacent

# Block Permuted Hamming Search

- Let,  $h=3$  and  $b=64$
- Take  $G=8$  blocks, so  $\gamma=64/8=8$
- Select integer  $g$  between 1 and  $G-h$ , say  $g=3$
- Comparison between  $x$  and  $y$  is guaranteed to have at least  $g$  exact block matches
- Use  $g$  blocks as header
- We can choose header in  $\binom{G}{g}$  ways. Each works as a permutation  $\pi_i$ .



# Conclusion

- Experiment on 1000 documents:
- Time taken by pairwise comparison: 625 sec
- Time taken by block permutation: 2.62 sec
- Number of pairs having similarity more than 90%: 19
- Detected pairs: 15



# References

- Charikar, Moses S. "Similarity estimation techniques from rounding algorithms." In *Proceedings of the thirty-fourth annual ACM symposium on Theory of computing*, pp. 380-388. ACM, 2002.
- Croft, W. Bruce, Donald Metzler, and Trevor Strohman. *Search engines: Information retrieval in practice*. Vol. 283. Reading: Addison-Wesley, 2010.
- Manku, Gurmeet Singh, Arvind Jain, and Anish Das Sarma. "Detecting near-duplicates for web crawling." In *Proceedings of the 16th international conference on World Wide Web*, pp. 141-150. ACM, 2007.
- Sood, Sadhan, and Dmitri Loguinov. "Probabilistic near-duplicate detection using simhash." In *Proceedings of the 20th ACM international conference on Information and knowledge management*, pp. 1117-1126. ACM, 2011.
- The Simhash Algorithm. Web: <http://matpalm.com/resemblance/simhash/>